

Vector Space Model

Slides borrowed from Hongning Wang with modification

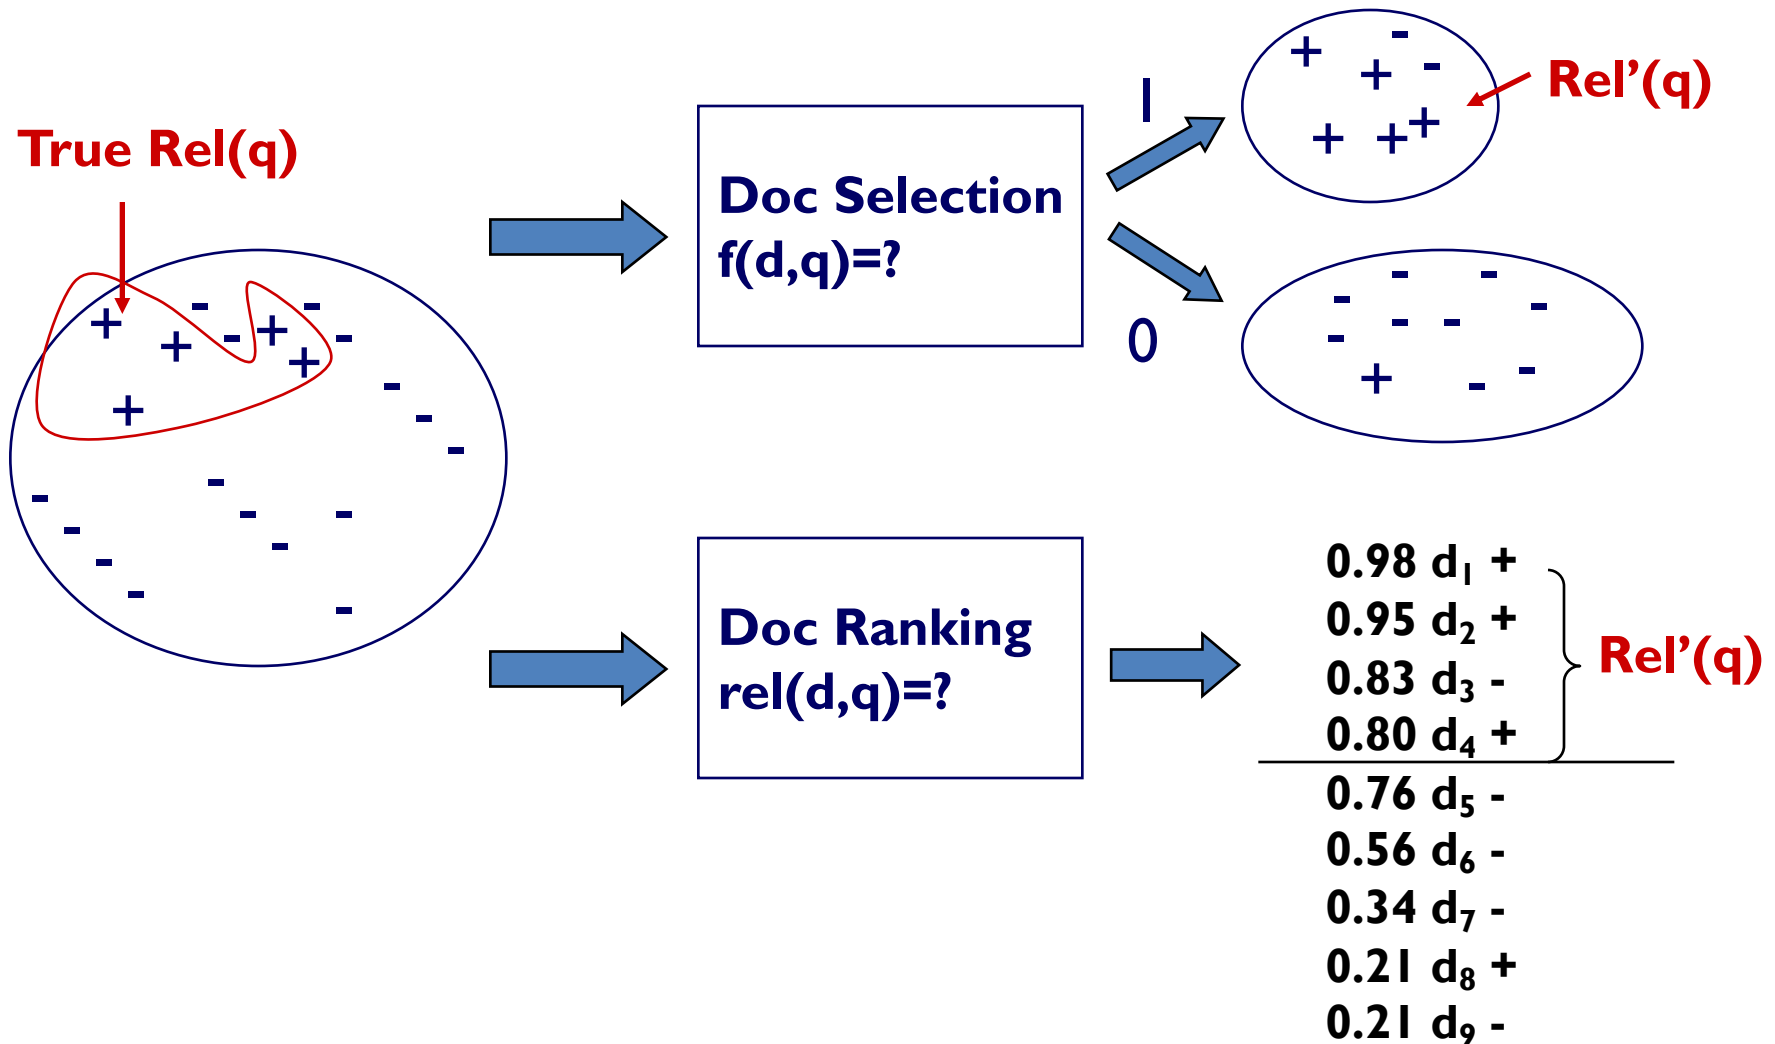
Ranked retrieval

- Thus far, our queries have all been Boolean
 - Documents either match or don't
- Can be good for expert users with precise understanding of their needs and the collection
 - Can also be good for applications: Applications can easily consume 1000s of results
- Not good for the majority of users
 - Most users incapable of writing Boolean queries
 - Or they are, but they think it's too much work
 - Most users don't want to wade through 1000s of results.
 - This is particularly true of web search

Problem with Boolean search: feast or famine

- Boolean queries often result in either too few (=0) or too many (1000s) results.
- Query 1: “*standard user dlink 650*” → 200,000 hits
- Query 2: “*standard user dlink 650 no card found*”: 0 hits
- It takes a lot of skill to come up with a query that produces a manageable number of hits.
 - AND gives too few; OR gives too many

Document Selection vs. Ranking



Intuitions for Ranking


- Query side: some terms are more important than others to represent the user's information need
- Document side: some terms carry more information about the document

Vector space model

- Represent both document and query by concept vectors
 - Each concept defines one dimension
 - K concepts define a high-dimensional space
 - Element of vector corresponds to concept weight
 - E.g., $d=(x_1, \dots, x_k)$, x_i is “importance” of concept i
- Measure relevance
 - Similarity between the query vector and document vector in this concept space

VS Model: an illustration

- Which document is closer to the query?




What the VS model doesn't say

- How to define/select the “basic concept”
 - Concepts are assumed to be orthogonal
- How to assign weights
 - Weight in query indicates importance of the concept
 - Weight in doc indicates how well the concept characterizes the doc
- How to define the similarity/distance measure

What is a good “basic concept”?

- Orthogonal
 - Linearly independent basis vectors
 - “Non-overlapping” in meaning
 - No ambiguity
- Weights can be assigned automatically and accurately
- Existing solutions
 - Terms or N-grams, i.e., bag-of-words
 - Topics, i.e., topic model

Bag of words representation



How to assign weights?

- Important!
- How?
 - Two basic heuristics
 - TF (Term Frequency) = Within-doc-frequency
 - IDF (Inverse Document Frequency)

TF weighting

- Idea: a term is more important if it occurs more frequently in a document
- TF Formulas
 - Let $f(t, d)$ be the frequency count of term t in doc d
 - Raw TF: $tf(t, d) = f(t, d)$


TF normalization

- Query: *iphone 6s*
 - D1: iPhone 6s receives pre-orders on September 12.
 - D2: iPhone 6 has three color options.
 - D3: iPhone 6 has three color options. iPhone 6 has three color options. iPhone 6 has three color options.

TF normalization

- Sublinear TF scaling

$$- \quad tf(t, d) = \begin{cases} 1 + \log f(t, d), & \text{if } f(t, d) > 0 \\ 0, & \text{otherwise} \end{cases}$$



Document frequency

- Idea: a term is more discriminative if it occurs only in fewer documents

IDF weighting

- Solution

- Assign higher weights to the rare terms

- Formula

- $IDF(t) = \log\left(\frac{N}{df(t)}\right)$

Non-linear scaling

Total number of docs in collection

Number of docs containing term t

- A corpus-specific property

- Independent of a single document

Collection vs. Document frequency

- Collection frequency of t is the total number of occurrences of t in the collection (incl. multiples)
- Document frequency is number of docs t is in
- Example:

Word	Collection frequency	Document frequency
<i>insurance</i>	10440	3997
<i>try</i>	10422	8760

- Which word is a better search term (and should get a higher weight)?

tf-idf weighting has many variants

Term frequency		Document frequency		Normalization	
n (natural)	$tf_{t,d}$	n (no)	1	n (none)	1
l (logarithm)	$1 + \log(tf_{t,d})$	t (idf)	$\log \frac{N}{df_t}$	c (cosine)	$\frac{1}{\sqrt{w_1^2 + w_2^2 + \dots + w_M^2}}$
a (augmented)	$0.5 + \frac{0.5 \times tf_{t,d}}{\max_t(tf_{t,d})}$	p (prob idf)	$\max\{0, \log \frac{N - df_t}{df_t}\}$	u (pivoted unique)	$1/u$
b (boolean)	$\begin{cases} 1 & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$			b (byte size)	$1/CharLength^\alpha$, $\alpha < 1$
L (log ave)	$\frac{1 + \log(tf_{t,d})}{1 + \log(\text{ave}_{t \in d}(tf_{t,d}))}$				

TF-IDF weighting

- Combining TF and IDF
 - Common in doc \rightarrow high tf \rightarrow high weight
 - Rare in collection \rightarrow high idf \rightarrow high weight
 - $w(t, d) = TF(t, d) \times IDF(t)$
- Most well-known document representation schema in IR! (G Salton et al. 1983)




"Salton was perhaps the leading computer scientist working in the field of information retrieval during his time." - wikipedia

[Gerard Salton Award](#)

– highest achievement award in IR

Cosine similarity

- Angle between two vectors
 - $\text{cosine}(V_q, V_d) = \frac{V_q \times V_d}{|V_q|_2 \times |V_d|_2} = \frac{V_q}{|V_q|_2} \times \frac{V_d}{|V_d|_2}$
 - Document length normalized
- TF-IDF vector
- Unit vector



What you should know

- Basic idea of vector space model
- Two important heuristics in VS model
 - TF
 - IDF
- Similarity measure for VS model
 - cosine similarity

Today's reading

- Chapter 6: Scoring, term weighting and the vector space model
 - 6.2 Term frequency and weighting
 - 6.3 The vector space model for scoring
 - 6.4 Variant tf-idf functions